

# Assignment 6.1

October 12, 2020

## 1 Assignment 6.1

Brandon Sams

05Oct2020

Using section 5.1 in Deep Learning with Python as a guide (listing 5.3 in particular), create a ConvNet model that classifies images in the MNIST digit dataset. Save the model, predictions, metrics, and validation plots in the dsc650/assignments/assignment06/results directory. If you are using JupyterHub, you can include those plots in your Jupyter notebook.

### 1.1 Instantiating a small convnet

```
[1]: from keras import layers
     from keras import models
```

```
[2]: model = models.Sequential()
     model.add(layers.Conv2D(32,(3,3),activation='relu',input_shape=(28,28,1)))
     model.add(layers.MaxPooling2D((2,2)))
     model.add(layers.Conv2D(64,(3,3),activation='relu',input_shape=(28,28,1)))
     model.add(layers.MaxPooling2D((2,2)))
     model.add(layers.Conv2D(64,(3,3),activation='relu',input_shape=(28,28,1)))
```

```
[3]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 64)	36928

```
Total params: 55,744
Trainable params: 55,744
Non-trainable params: 0
```

---

## 1.2 Adding a classifier on top of the convnet

```
[4]: model.add(layers.Flatten())
      model.add(layers.Dense(64,activation='relu'))
      model.add(layers.Dense(10,activation='softmax'))
```

```
[5]: model.summary()
```

```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 64)	0
conv2d_2 (Conv2D)	(None, 3, 3, 64)	36928
flatten (Flatten)	(None, 576)	0
dense (Dense)	(None, 64)	36928
dense_1 (Dense)	(None, 10)	650

```
Total params: 93,322
Trainable params: 93,322
Non-trainable params: 0
```

---

## 1.3 Training the convnet on MNIST images

```
[6]: from keras.datasets import mnist
      from keras.utils import to_categorical
```

```
[7]: (train_images, train_labels), (test_images, test_labels) = mnist.load_data()

      train_images = train_images.reshape((60000, 28, 28, 1))
      train_images = train_images.astype('float32') / 255
```

```
test_images = test_images.reshape((10000, 28, 28, 1))
test_images = test_images.astype('float32') / 255

train_labels = to_categorical(train_labels)
test_labels = to_categorical(test_labels)
```

```
[8]: model.compile(optimizer='rmsprop',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])

model.fit(train_images,
          train_labels,
          epochs=50,
          batch_size=64)
```

```
Epoch 1/50
938/938 [=====] - 15s 16ms/step - loss: 0.1699 -
accuracy: 0.9464
Epoch 2/50
938/938 [=====] - 15s 16ms/step - loss: 0.0480 -
accuracy: 0.9853
Epoch 3/50
938/938 [=====] - 15s 15ms/step - loss: 0.0329 -
accuracy: 0.9900
Epoch 4/50
938/938 [=====] - 14s 15ms/step - loss: 0.0246 -
accuracy: 0.9923
Epoch 5/50
938/938 [=====] - 14s 15ms/step - loss: 0.0193 -
accuracy: 0.9943
Epoch 6/50
938/938 [=====] - 14s 15ms/step - loss: 0.0154 -
accuracy: 0.9952
Epoch 7/50
938/938 [=====] - 14s 15ms/step - loss: 0.0118 -
accuracy: 0.9964
Epoch 8/50
938/938 [=====] - 14s 15ms/step - loss: 0.0107 -
accuracy: 0.9967
Epoch 9/50
938/938 [=====] - 14s 15ms/step - loss: 0.0089 -
accuracy: 0.9972
Epoch 10/50
938/938 [=====] - 14s 15ms/step - loss: 0.0071 -
accuracy: 0.9978
Epoch 11/50
```

938/938 [=====] - 14s 15ms/step - loss: 0.0063 -  
 accuracy: 0.9981  
 Epoch 12/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0052 -  
 accuracy: 0.9986  
 Epoch 13/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0052 -  
 accuracy: 0.9985  
 Epoch 14/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0046 -  
 accuracy: 0.9987  
 Epoch 15/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0052 -  
 accuracy: 0.9986  
 Epoch 16/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0036 -  
 accuracy: 0.9990  
 Epoch 17/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0033 -  
 accuracy: 0.9991  
 Epoch 18/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0030 -  
 accuracy: 0.9991  
 Epoch 19/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0033 -  
 accuracy: 0.9993  
 Epoch 20/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0028 -  
 accuracy: 0.9992  
 Epoch 21/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0034 -  
 accuracy: 0.9991  
 Epoch 22/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0041 -  
 accuracy: 0.9991  
 Epoch 23/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0026 -  
 accuracy: 0.9993  
 Epoch 24/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0028 -  
 accuracy: 0.9993  
 Epoch 25/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0032 -  
 accuracy: 0.9993  
 Epoch 26/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0034 -  
 accuracy: 0.9991  
 Epoch 27/50

938/938 [=====] - 14s 15ms/step - loss: 0.0021 -  
 accuracy: 0.9995  
 Epoch 28/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0012 -  
 accuracy: 0.9997  
 Epoch 29/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0019 -  
 accuracy: 0.9995  
 Epoch 30/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0027 -  
 accuracy: 0.9993  
 Epoch 31/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0026 -  
 accuracy: 0.9993  
 Epoch 32/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0018 -  
 accuracy: 0.9995  
 Epoch 33/50  
 938/938 [=====] - 15s 16ms/step - loss: 0.0019 -  
 accuracy: 0.9995  
 Epoch 34/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0025 -  
 accuracy: 0.9995  
 Epoch 35/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0011 -  
 accuracy: 0.9997  
 Epoch 36/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0015 -  
 accuracy: 0.9997  
 Epoch 37/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0014 -  
 accuracy: 0.9997  
 Epoch 38/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0013 -  
 accuracy: 0.9997  
 Epoch 39/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0016 -  
 accuracy: 0.9997  
 Epoch 40/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0019 -  
 accuracy: 0.9997  
 Epoch 41/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0019 -  
 accuracy: 0.9995  
 Epoch 42/50  
 938/938 [=====] - 14s 15ms/step - loss: 0.0010 -  
 accuracy: 0.9998  
 Epoch 43/50

```

938/938 [=====] - 14s 15ms/step - loss: 0.0031 -
accuracy: 0.9995
Epoch 44/50
938/938 [=====] - 14s 15ms/step - loss: 0.0016 -
accuracy: 0.9997
Epoch 45/50
938/938 [=====] - 14s 15ms/step - loss: 0.0025 -
accuracy: 0.9996
Epoch 46/50
938/938 [=====] - 14s 15ms/step - loss: 8.1034e-04 -
accuracy: 0.9998
Epoch 47/50
938/938 [=====] - 14s 15ms/step - loss: 0.0023 -
accuracy: 0.9995
Epoch 48/50
938/938 [=====] - 14s 15ms/step - loss: 0.0025 -
accuracy: 0.9996
Epoch 49/50
938/938 [=====] - 14s 15ms/step - loss: 0.0024 -
accuracy: 0.9996
Epoch 50/50
938/938 [=====] - 14s 15ms/step - loss: 0.0033 -
accuracy: 0.9994

```

[8]: <tensorflow.python.keras.callbacks.History at 0x7f456ed90760>

```

[9]: # Let's evaluate the model on the test data
test_loss, test_acc = model.evaluate(test_images, test_labels)
test_acc

```

```

313/313 [=====] - 1s 5ms/step - loss: 0.1693 -
accuracy: 0.9907

```

[9]: 0.9907000064849854

## 1.4 Save model

```

[10]: model.save('./results/model_6_1.h5')

```

```

[11]: import pandas as pd
import numpy as np

num_testing = len(test_images)
pred = np.argmax(model.predict(test_images), axis=-1)
sub_df = pd.DataFrame()
sub_df["ImageId"] = list(range(1, num_testing + 1))
sub_df["Label"] = pred

```

```
[12]: (train_images, train_labels), (test_images, test_labels) = mnist.load_data()

sub_df["True_Label"] = test_labels
```

```
[13]: sub_df
```

```
[13]:
```

	ImageId	Label	True_Label
0	1	7	7
1	2	2	2
2	3	1	1
3	4	0	0
4	5	4	4
...	...	...	...
9995	9996	2	2
9996	9997	3	3
9997	9998	4	4
9998	9999	5	5
9999	10000	6	6

[10000 rows x 3 columns]

```
[14]: sub_df[sub_df.Label != sub_df.True_Label]
```

```
[14]:
```

	ImageId	Label	True_Label
62	63	5	9
104	105	5	9
321	322	7	2
340	341	3	5
582	583	2	8
...	...	...	...
9664	9665	7	2
9698	9699	1	6
9700	9701	4	2
9729	9730	6	5
9839	9840	7	2

[93 rows x 3 columns]

```
[ ]:
```